

## Fraud Analytics Assignment

# Title: Example-dependent cost-sensitive regression

### Group members:

**Name : Dindorkar Mayuresh Rajesh**

**Roll No : CS23MTECH14007**

**Name : Sanyam Kaul**

**Roll No : CS23MTECH14011**

**Name : Shrenik Ganguli**

**Roll No : CS23MTECH14014**

**Name : Morey Piyush Prabhakar**

**Roll No : AI23MTECH14003**

```
In [1]: # Library used for genetic algorithm
!pip install pygad
```

```
Requirement already satisfied: pygad in c:\users\sanya\anaconda3\lib\site-packages (3.3.1)
Requirement already satisfied: cloudpickle in c:\users\sanya\anaconda3\lib\site-packages (from pygad) (2.2.1)
Requirement already satisfied: matplotlib in c:\users\sanya\anaconda3\lib\site-packages (from pygad) (3.7.2)
Requirement already satisfied: numpy in c:\users\sanya\anaconda3\lib\site-packages (from pygad) (1.24.3)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (1.0.5)
Requirement already satisfied: cyclers>=0.10 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (23.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (9.4.0)
Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\sanya\anaconda3\lib\site-packages (from matplotlib->pygad) (2.8.2)
Requirement already satisfied: six>=1.5 in c:\users\sanya\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib->pygad) (1.16.0)
```

### Importing libraries:

```
In [2]: import numpy as np
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.preprocessing import StandardScaler
import warnings
```

```
from sklearn.linear_model import LogisticRegression
import seaborn as sns
warnings.filterwarnings("ignore")
import pygad
from scipy.optimize import minimize
```

### Utility Functions:

```
In [3]: def count_labels(data, column_name):
        return data[column_name].value_counts()
```

```
In [4]: def plot_bar_chart(label_counts):
        labels = label_counts.index.tolist()
        counts = label_counts.values
        plt.bar(np.arange(len(labels)), counts, color=['blue', 'orange'])
        plt.ylabel('Count')
        plt.xlabel('Status')
        plt.title(f'Count of labels in the Status column')
        plt.xticks(np.arange(len(labels)), labels)
        plt.show()
```

```
In [5]: def perform_scaling(X_train, X_test):
        std_scaler = StandardScaler()
        X_train = std_scaler.fit_transform(X_train)
        X_test = std_scaler.transform(X_test)
        return std_scaler, X_train, X_test
```

```
In [6]: # Returns FP, TP, TN, FNC
        def get_cost_matrix_values(Y_train):
            num_of_samples = X_train.shape[0]
            return np.full(num_of_samples, 6), np.full(num_of_samples, 6), np.zeros(num_of_
```

```
In [7]: def compute_sigmoid(a):
        return 1 / (1 + np.exp(-a))
```

```
In [8]: def calculate_saving(simple_lr_cost, cost_of_csir):
        return (simple_lr_cost - cost_of_csir) / simple_lr_cost
```

### Reading the dataset:

```
In [9]: df = pd.read_csv('costsensitiveregession.csv')
        df.head()
```

```
Out[9]:
```

|   | NotCount | YesCount | ATPM | PFD   | PFG | SFD | SFG | WP       | WS  | AH  | AN  | Status | FNC |
|---|----------|----------|------|-------|-----|-----|-----|----------|-----|-----|-----|--------|-----|
| 0 | 2        | 21       | 0.0  | 0.000 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0      | 0.0 |
| 1 | 23       | 0        | 0.0  | 0.044 | 0.0 | 0.0 | 0.0 | 0.306179 | 0.0 | 0.0 | 0.0 | 1      | 0.0 |
| 2 | 1        | 22       | 0.0  | 0.000 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0      | 0.0 |
| 3 | 5        | 18       | 0.0  | 0.000 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 1      | 0.0 |
| 4 | 1        | 22       | 0.0  | 0.000 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0      | 0.0 |

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 147636 entries, 0 to 147635
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   NotCount    147636 non-null  int64
1   YesCount    147636 non-null  int64
2   ATPM        147636 non-null  float64
3   PFD         147636 non-null  float64
4   PFG         147636 non-null  float64
5   SFD         147636 non-null  float64
6   SFG         147636 non-null  float64
7   WP          147636 non-null  float64
8   WS          147636 non-null  float64
9   AH          147636 non-null  float64
10  AN          147636 non-null  float64
11  Status      147636 non-null  int64
12  FNC         147636 non-null  float64
dtypes: float64(10), int64(3)
memory usage: 14.6 MB
```

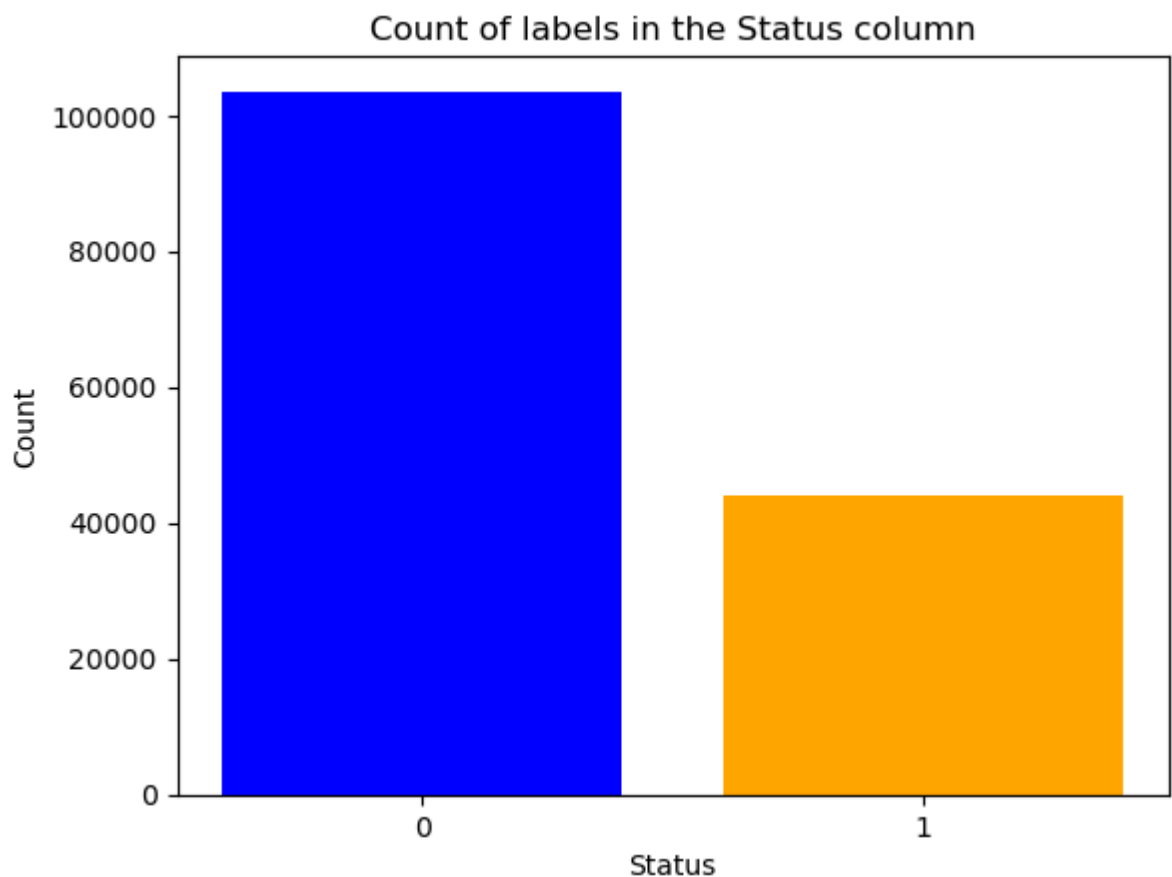
### Checking null values in the dataset:

```
In [11]: df.isnull().sum()
```

```
Out[11]: NotCount    0
YesCount    0
ATPM        0
PFD         0
PFG         0
SFD         0
SFG         0
WP          0
WS          0
AH          0
AN          0
Status      0
FNC         0
dtype: int64
```

### Count plot for feature 'Status':

```
In [12]: label_counts = count_labels(df, 'Status')
plot_bar_chart(label_counts)
```



### Data Pre-processing:

#### Segregating the dataset into dependant and Independent features:

```
In [13]: X, Y = df.iloc[:, :-2], df.iloc[:, -2:]
```

#### Segregating dataset into train and test sets:

```
In [14]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

#### Scaling the data

```
In [15]: std_scaler, X_train, X_test = perform_scaling(X_train, X_test)
```

## Bahnsen's Approach:

#### Creating cost matrix:

```
In [16]: FP, TP, TN, FNC = get_cost_matrix_values(Y_train)
```

#### Cost Sensitive Logistic regression Implementation:

$$J^c(\theta) = \frac{1}{N} \sum_{i=1}^N \left( y_i(h_{\theta}(\mathbf{x}_i)C_{TP_i} + (1 - h_{\theta}(\mathbf{x}_i))C_{FN_i}) + (1 - y_i)(h_{\theta}(\mathbf{x}_i)C_{FP_i} + (1 - h_{\theta}(\mathbf{x}_i))C_{TN_i}) \right).$$

```
In [17]: TO_AVOID_DIVIDE_BY_ZERO_ERROR = 1e-08
def cost_sensitive_logistic_regression_cost_function(Y, Y_hat, TP, TN, FP, FN):
    cost = (1 - Y) * (Y_hat * FP + (1 - Y_hat) * TN)
    return np.mean(cost + (Y * (Y_hat * TP + (1 - Y_hat) * FN)))

def calculate_fitness(generation_instance, weights, weight_index):
    Y_hat = compute_sigmoid(X_train @ weights)
    return np.reciprocal(cost_sensitive_logistic_regression_cost_function(Y = Y_hat, Y_hat = Y_hat, TP = TP, TN = TN, FP = FP, FN = FN))

class ExampleDependantCostSensitiveRegression():
    def __init__(self, generation_iterations = 20, count_parents_mating = 6, training_function = train, type_of_parent_selection = 'sss', per_pop_solutions = 35, number_of_genes = 10, per_gene_mutation_percent = 10, type_of_mutation = 'random', printing_function = lambda ga_instance: print('Generation: {} have fitness: {}'.format(ga_instance.generation, ga_instance.best_solution().fitness))):
        super(ExampleDependantCostSensitiveRegression, self).__init__()
        self.type_of_crossover = type_of_crossover
        self.count_parents_mating = count_parents_mating
        self.printing_function = printing_function
        self.training_function = training_function
        self.per_pop_solutions = per_pop_solutions
        self.number_of_genes = number_of_genes
        self.generation_iterations = generation_iterations
        self.type_of_parent_selection = type_of_parent_selection
        self.per_gene_mutation_percent = per_gene_mutation_percent
        self.keep_parents = keep_parents
        self.type_of_mutation = type_of_mutation

    def get_genetic_instance(self):
        return pygad.GA(fitness_func = self.training_function, mutation_type = self.type_of_mutation, parent_selection_type = self.type_of_parent_selection, solution_validation = self.printing_function, num_genes = self.number_of_genes, mutation_percent_genes = self.per_gene_mutation_percent)

    def train(self, genetic_instance):
        genetic_instance.run()

    def get_trained_weights(self, genetic_instance):
        return genetic_instance.best_solution()

    def predict(self, X_test, weights):
        print(X_test.shape)
        return compute_sigmoid(X_test @ weights)
```

### Runnung the genetic algorithm to obtain best weights of cost sensitive logistic regression:

```
In [18]: obj = ExampleDependantCostSensitiveRegression()
genetic_instance = obj.get_genetic_instance()
obj.train(genetic_instance)
```

Generation: 1 have Fitness value: 0.1785177923840861  
 Generation: 2 have Fitness value: 0.23552389919258399  
 Generation: 3 have Fitness value: 0.23797253904953258  
 Generation: 4 have Fitness value: 0.24101333065378686  
 Generation: 5 have Fitness value: 0.24367988868208634  
 Generation: 6 have Fitness value: 0.24725999218699146  
 Generation: 7 have Fitness value: 0.247991582409019  
 Generation: 8 have Fitness value: 0.25082783431455935  
 Generation: 9 have Fitness value: 0.2521244468433215  
 Generation: 10 have Fitness value: 0.2527666828003852  
 Generation: 11 have Fitness value: 0.25375708711372835  
 Generation: 12 have Fitness value: 0.2545431708755377  
 Generation: 13 have Fitness value: 0.2564385236785665  
 Generation: 14 have Fitness value: 0.25750940827977065  
 Generation: 15 have Fitness value: 0.2577795767401475  
 Generation: 16 have Fitness value: 0.25864999669636574  
 Generation: 17 have Fitness value: 0.2591323707833581  
 Generation: 18 have Fitness value: 0.260169090703135  
 Generation: 19 have Fitness value: 0.26024322265451755  
 Generation: 20 have Fitness value: 0.2605248751790192

```

In [19]: trained_weights, _, _ = obj.get_trained_weights(genetic_instance)
         print('Trained weights:', trained_weights)
         predictions = obj.predict(X_test, trained_weights)
         predictions = np.array([1 if x >= 0.5 else 0 for x in predictions])
         cost_of_csrlr = cost_sensitive_logistic_regression_cost_function(Y = Y_test.iloc[:,0], predictions = predictions)
         print("Cost for Cost Sensitive Logistic Regression:", cost_of_csrlr)
  
```

Trained weights: [ 5.79669059 -7.0959525 5.99533002 3.19614711 -0.97488652 2.37395127  
 0.29067918 3.14980413 2.49213836 0.32127434 0.0661056 ]  
 (29528, 11)  
 Cost for Cost Sensitive Logistic Regression: 3.8526688022216202

---

### Comparing the results with inbuilt Logistic Regressor of ScikitLearn:

```

In [20]: simple_lr = LogisticRegression()
         simple_lr.fit(X_train, Y_train.iloc[:,0])
         simple_lr_pred = simple_lr.predict(X_test)

         simple_lr_pred = np.array([1 if x >= 0.5 else 0 for x in simple_lr_pred])
         simple_lr_cost = cost_sensitive_logistic_regression_cost_function(Y = Y_test.iloc[:,0], predictions = simple_lr_pred)
         print("Cost for Simple (Sklearn's) Logistic Regression:", simple_lr_cost)
  
```

Cost for Simple (Sklearn's) Logistic Regression: 27.71167158971146

---

### Calculating the cost saving score:

```

In [21]: print("Simple LR score Vs Cost sensitive LR saving score:", calculate_saving(cost_of_csrlr, simple_lr_cost))
         Simple LR score Vs Cost sensitive LR saving score: 0.8609730636512016
  
```

---

## Guhnmann's Approach (Variant A):

$$\int_0^1 a_i \cdot (y_i \cdot (-\log f(g(x_i, \beta)))) df \stackrel{!}{=} c_i$$

```
In [22]: def weighted_logistic_loss(weights, X, y, cost_matrix):
# Extract model parameters
w = weights[:-1]
b = weights[-1]

# Compute predicted probabilities
z = np.dot(X, w) + b
probs = 1 / (1 + np.exp(-z))

# Compute Logistic Loss with misclassification costs
loss = -np.mean(cost_matrix[:, 0] * y[:,0] * np.log(probs + 1e-15 ) +
               cost_matrix[:, 1] * y[:,0] * np.log(1 - probs + 1e-15 ) +
               cost_matrix[:, 2] * (1 - y[:,0]) * np.log(probs + 1e-15 ) +
               cost_matrix[:, 3] * (1 - y[:,0]) * np.log(1 - probs + 1e-15 ))

return loss
```

```
In [23]: def cost_sensitive_logistic_regression(X_train, y_train, X_test, y_test, cost_matrix):
# Initialize weights (including bias term)
num_features = X_train.shape[1]
initial_weights = np.zeros(num_features + 1)

# Define callback function to print iteration information
def callback(weights):
    print(f"Loss: {weighted_logistic_loss(weights, X_train, y_train, cost_matrix)}")

# Minimize the weighted logistic loss with iteration callback
res = minimize(weighted_logistic_loss, initial_weights, args=(X_train, y_train, cost_matrix))
weights = res.x

# Extract optimal weights and bias
w = weights[:-1]
b = weights[-1]

# Predict on test data
z = np.dot(X_test, w) + b
y_pred = np.round(1 / (1 + np.exp(-z)))

print('Finished')
return weights
```

```
In [24]: # Split features and target variable
X = df.iloc[:, :-2].values # Features (columns 0 to 10)
y = df.iloc[:, -2:].values # Target variable (column 11)

# Standardize features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define the cost matrix based on your problem
False_positive = np.full(len(y_train), 4)
True_positive = np.full(len(y_train), 4)
True_negative = np.zeros(len(y_train))
cost_matrix = np.column_stack((False_positive, y_train[:,1], True_positive, True_negative))
```

```
# Train cost-sensitive Logistic regression model
weights = cost_sensitive_logistic_regression(X_train, y_train, X_test, y_test, cost_matrix)

Loss: 18.285193352358363
Loss: 12.109292869353467
Loss: 11.605889248016902
Loss: 10.854695121334833
Loss: 10.005322628099108
Loss: 9.938750775096162
Loss: 9.82102061061701
Loss: 9.654153797429135
Loss: 9.413174322729056
Loss: 9.159201120332899
Loss: 8.853975828469498
Loss: 8.630809870690786
Loss: 8.535067006824622
Loss: 8.494091806095675
Loss: 8.477056068355825
Loss: 8.469877809773024
Loss: 8.466031656435044
Loss: 8.462488766571093
Loss: 8.459722469851604
Loss: 8.457842755576396
Finished
```

```
In [25]: cost_of_csrlr = weighted_logistic_loss(weights,X_train, y_train, cost_matrix)
```

```
In [26]: print("Cost for Cost Sensitive Logistic Regression (Guhnmann's approach):", cost_of_csrlr)

Cost for Cost Sensitive Logistic Regression (Guhnmann's approach): 8.457842755576396
```

```
In [27]: print("Simple LR score Vs Cost sensitive LR saving (Guhnmann's approach) score:", cost_of_csrlr)

Simple LR score Vs Cost sensitive LR saving (Guhnmann's approach) score: 0.6947913182286504
```

Results:

- Ran all the logistic regression algorithms for 20 epochs

| Approach  | Average Cost | Savings Score wtr Simple Logistic Regression |
|---|--------------|--|
| Simple Logistic Regression                              | 27.71        | NA   |
| Cost Sensitive Logistic Regression (Bahnsen's Approach) | 3.93         | 0.85   |
| Cost Sensitive Logistic Regression (Guhnman's Approach) | 8.45         | 0.69   |

```
In [ ]:
```